Лекция 3

Инструменты профилирования:

- nvprof
- Nsight Compute CLI
- nvvp
- Nsight Compute

nvprof и Nsight Compute CLI

```
__global__ void gInit(float* a, float* b){
  int i=threadIdx.x+blockIdx.x*blockDim.x;
  a[i]=(float)2*i;
  b[i]=(float)(2*i+1);
}

__global__ void gSum(float* a, float *b){
  int i=threadIdx.x+blockIdx.x*blockDim.x;
  a[i]+=b[i];
}
```

Тестовые ядра

```
ip-011@linux-47dw:/home/malkov/WORKSHOP/PGP-2023>
nvprof ./lab3c
```

```
Type  Time(%) Time Calls Avg Min Max Name
GPU activities:
43.59%  2.1760us  1  2.1760us  2.1760us
                                  2.1760us  gSum(int*, int*)
41.67%  2.0800us  1  2.0800us  2.0800us
                                   2.0800us  gInit(int*, int*)
14.74%     736ns  1  736ns   736ns   736ns  [CUDA memcpy DtoH]
API calls:
98.87%  131.54ms  2  65.772ms  6.9650us  131.54ms  cudaMalloc
……..
0.09%  124.46us   2  62.229us  10.561us  113.90us  cudaFree
……
0.01%  14.599us   1  14.599us  14.599us  14.599us

                                        cudaMemcpy
```

```
/Lecture3/Lab3-cuda-gdb # ncu --target-processes all  ./lab3c

 gInit(int *, int *), 2023-Feb-13 15:09:06, Context 1, Stream 7
   Section: GPU Speed Of Light Throughput
---------------------------------------------------------------
DRAM Frequency cycle/nsecond                    6.40
SM Frequency   cycle/nsecond                    1.29
Elapsed Cycles         cycle                    3,327
Memory [%]                 %                    1.10
DRAM Throughput            %                    0.02
Duration             usecond                    2.56
 --------------- ------------------------------
   WRN   This kernel grid is too small to fill the available
resources on this device, resulting in only 0.0 full
waves across all SMs. Look at Launch Statistics for more
details.
```

……

```
/Lecture3/Lab3-cuda-gdb # ncu
--metrics gpu__time_duration.sum  ./lab3c
```

```
gInit(int *, int *), 2023-Feb-13 18:42:52, Context 1, Stream 7
   Section: Command line profiler metrics
   ------ --------------- ------------------------------
gpu__time_duration.sum   usecond                     29.50
----------------------------------------------------------
gSum(int *, int *), 2023-Feb-13 18:42:52, Context 1, Stream 7
   Section: Command line profiler metrics
----------------------------------------------------------
gpu__time_duration.sum   usecond                     37.57
----------------------------------------------------------
```

```
/Lecture3/Lab3-cuda-gdb> nvprof --query-metrics
======== Warning: Skipping profiling on device 0 since
profiling is not supported on devices with compute
capability 7.5 and higher.
      Use NVIDIA Nsight Compute for GPU profiling and NVIDIA
Nsight Systems for GPU tracing and CPU sampling.
        Refer https://developer.nvidia.com/tools-overview
for more details.
```

ip-011@linux-47dw:/home/malkov/WORKSHOP/PGP-2023>
nvprof --query-metrics | less

```
Available Metrics:    Name                    Description
Device 0 (GeForce GTX 1050):
inst_per_warp:  Average number of instructions executed by each
warp

warp_execution_efficiency:  Ratio of the average active threads
per warp to the maximum number of
threads per warp supported on a multiprocessor
……….
gld_transactions_per_request:  Average number of global memory
load transactions performed for each global memory load.

gst_transactions_per_request:  Average number of global memory
store transactions performed for each global memory store
……
```

```
ip-011@linux-47dw:/home/malkov/WORKSHOP/PGP-2023>
nvprof -m gst_throughput ./lab3c
```

```
Invocations Metric Name  Metric Description  Min   Max   Avg
Device "GeForce GTX 1050 (0)"
   Kernel: gSum(int*, int*)
 1       gst_throughput  Global Store Throughput  40.582MB/s
                                                 40.582MB/s  40.582MB/s
   Kernel: gInit(int*, int*)
 1       gst_throughput  Global Store Throughput  71.303MB/s
                                                  71.303MB/s  71.302MB/s
```

/Лекция4/lab4> **ncu --list-sections**

Identifier Display Name Enabled Filename

ComputeWorkloadAnalysis Compute Workload Analysis yes ...2024.2.1/Sections/ComputeWorkloadAnalysis.section

InstructionStats Instruction Statistics yes ...2024.2.1/Sections/InstructionStatistics.section

LaunchStats Launch Statistics yes ...2024.2.1/Sections/LaunchStatistics.section

MemoryWorkloadAnalysis Memory Workload Analysis yes …

……

/Лекция4/lab4> **ncu --section InstructionStats ./lab4c**

```
 gInit(float *, float *) (2, 1, 1)x(128, 1, 1), Context 1, Stream 7, Device 0, CC 7.5
   Section: Instruction Statistics
   ---------------------------------------- ----------- ------------
   Metric Name                              Metric Unit Metric Value
   ---------------------------------------- ----------- ------------
   Avg. Executed Instructions Per Scheduler        inst         0,93
   Executed Instructions                           inst          112
   Avg. Issued Instructions Per Scheduler          inst         1,27
   Issued Instructions                             inst          152
   ---------------------------------------- ----------- ------------
gSum(float *, float *) (2, 1, 1)x(128, 1, 1), Context 1,
Stream 7, Device 0, CC 7.5
   Section: Instruction Statistics
_----------------------------------------------------------------
```

/Лекция4/lab4> **ncu --section ComputeWorkloadAnalysis ./lab4c**
gSum(float *, float *) (2, 1, 1)x(128, 1, 1), Context 1, Stream 7, Device 0, CC 7.5
Section: Compute Workload Analysis
---------- ---------- ----------
Metric Name Metric Unit Metric Value
---------- ---------- ----------
Executed Ipc Active inst/cycle 0,04
Executed Ipc Elapsed inst/cycle 0,00
Issue Slots Busy % 1,35
Issued Ipc Active inst/cycle 0,05
SM Busy % 1,35
---------- ---------- ----------

OPT Est. Local Speedup: 99.33%
All compute pipelines are under-utilized. Either this kernel is very small or it doesn't issue enough warps per scheduler. Check the Launch Statistics and Scheduler Statistics sections for further details.

/Лекция4/lab4> **ncu --query-metrics > metrics.txt**

Device NVIDIA GeForce RTX 2060 (TU104)

Metric Name	Metric Type	Metric Unit	Metric Description
dram__bytes	Counter	byte	# of bytes accessed in DRAM
dram__bytes_read	Counter	byte	# of bytes read from DRAM
dram__bytes_write	Counter	byte	# of bytes written to DRAM

……..

smsp__average_inst_executed_pipe_lsu_per_warp Ratio inst/warp

average # of instructions executed by pipe lsu per warp

……..

/Лекция4/lab4> ***ncu --metrics l1tex__t_bytes_pipe_lsu_mem_global_op_st.sum.per_second ./lab4c***

```
gInit(float *, float *) (2, 1, 1)x(128, 1, 1), Context 1, Stream 7, Device 0, CC 7.5
  Section: Command line profiler metrics
  ---------------------------------------------------------- ----------- ------------
  Metric Name                                                Metric Unit Metric Value
  ---------------------------------------------------------- ----------- ------------
  l1tex__t_bytes_pipe_lsu_mem_global_op_st.sum.per_second     Mbyte/s       688,17
  ---------------------------------------------------------- ----------- ------------

gSum(float *, float *) (2, 1, 1)x(128, 1, 1), Context 1, Stream 7, Device 0, CC 7.5
  Section: Command line profiler metrics
  ---------------------------------------------------------- ----------- ------------
  Metric Name                                                Metric Unit Metric Value
  ---------------------------------------------------------- ----------- ------------
  l1tex__t_bytes_pipe_lsu_mem_global_op_st.sum.per_second     Mbyte/s       347,83
  ---------------------------------------------------------- ----------- ------------
```

Кодирование метрики ncu:

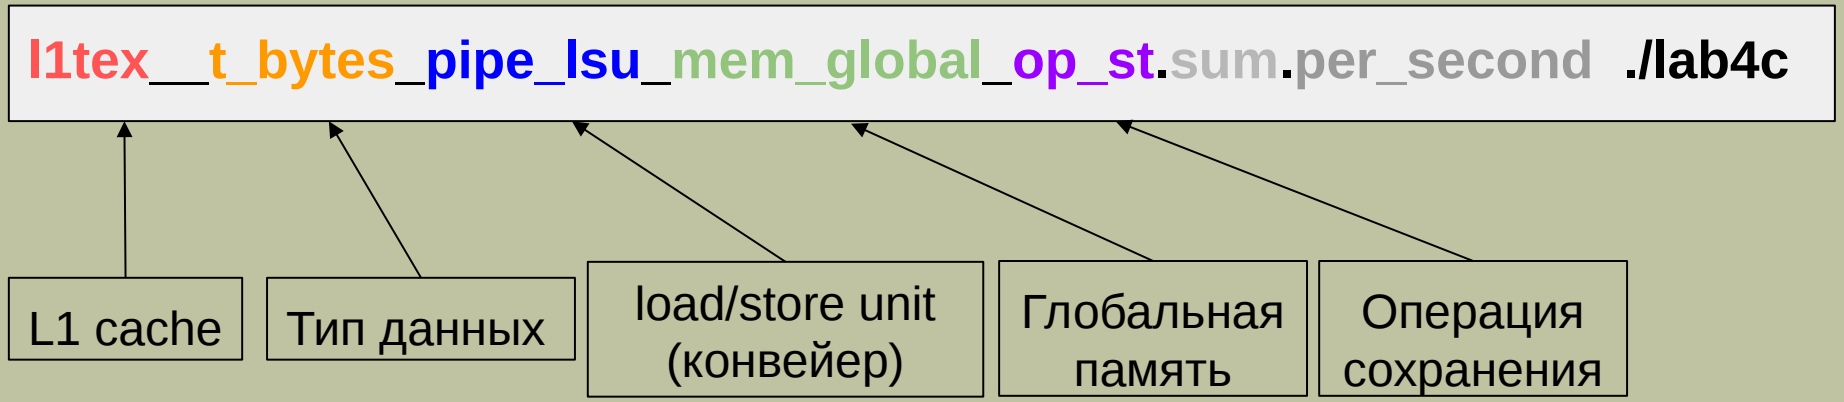

```
ip-011@linux-47dw:/home/malkov/WORKSHOP/PGP-2023>
nvprof -m gld_throughput ./lab3c
```

```
Invocations   Metric Name  Metric Description  Min Max   Avg
Device "GeForce GTX 1050 (0)"
   Kernel: gInit(int*, int*)
 1   gld_throughput Global Load Throughput 0.0B/s 0.0B/s  0.0B/s
   Kernel: gSum(int*, int*)
 1   gld_throughput Global Load Throughput  87.694MB/s

                    87.694MB/s  87.694MB/s
```

```
/Lecture3/Lab3-cuda-gdb # ncu --metrics
l1tex__t_bytes_pipe_lsu_mem_global_op_ld.sum.per_second
./lab3c
```

```
gInit(int *, int *), 2023-Feb-13 15:25:41, Context 1, Stream 7
   Section: Command line profiler metrics
 ---------------------------------------------------
l1tex__t_bytes_pipe_lsu_mem_global_op_ld.sum.per_second
         byte/second                                0
   ---------------------------------------------------
gSum(int *, int *), 2023-Feb-13 15:25:41, Context 1, Stream 7
   Section: Command line profiler metrics
   ---------------------------------------------------
l1tex__t_bytes_pipe_lsu_mem_global_op_ld.sum.per_second
        Mbyte/second                            82.47
   ---------------------------------------------------
```

nvvp и Nsight Compute

```
ip-011@linux-47dw:/home/malkov/WORKSHOP/PGP-2023> nvvp ./lab3c
```

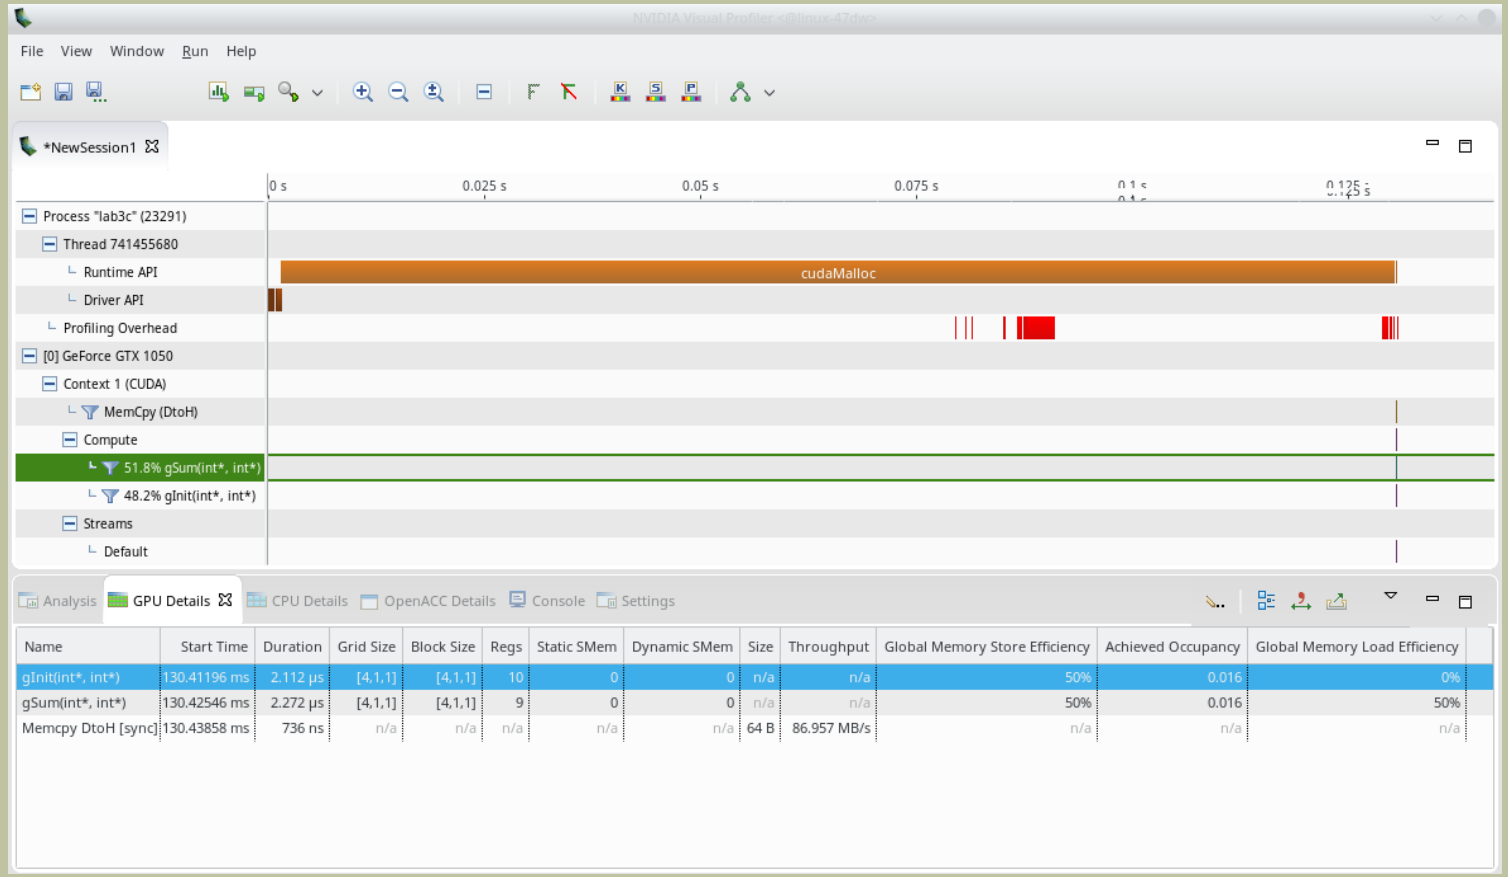

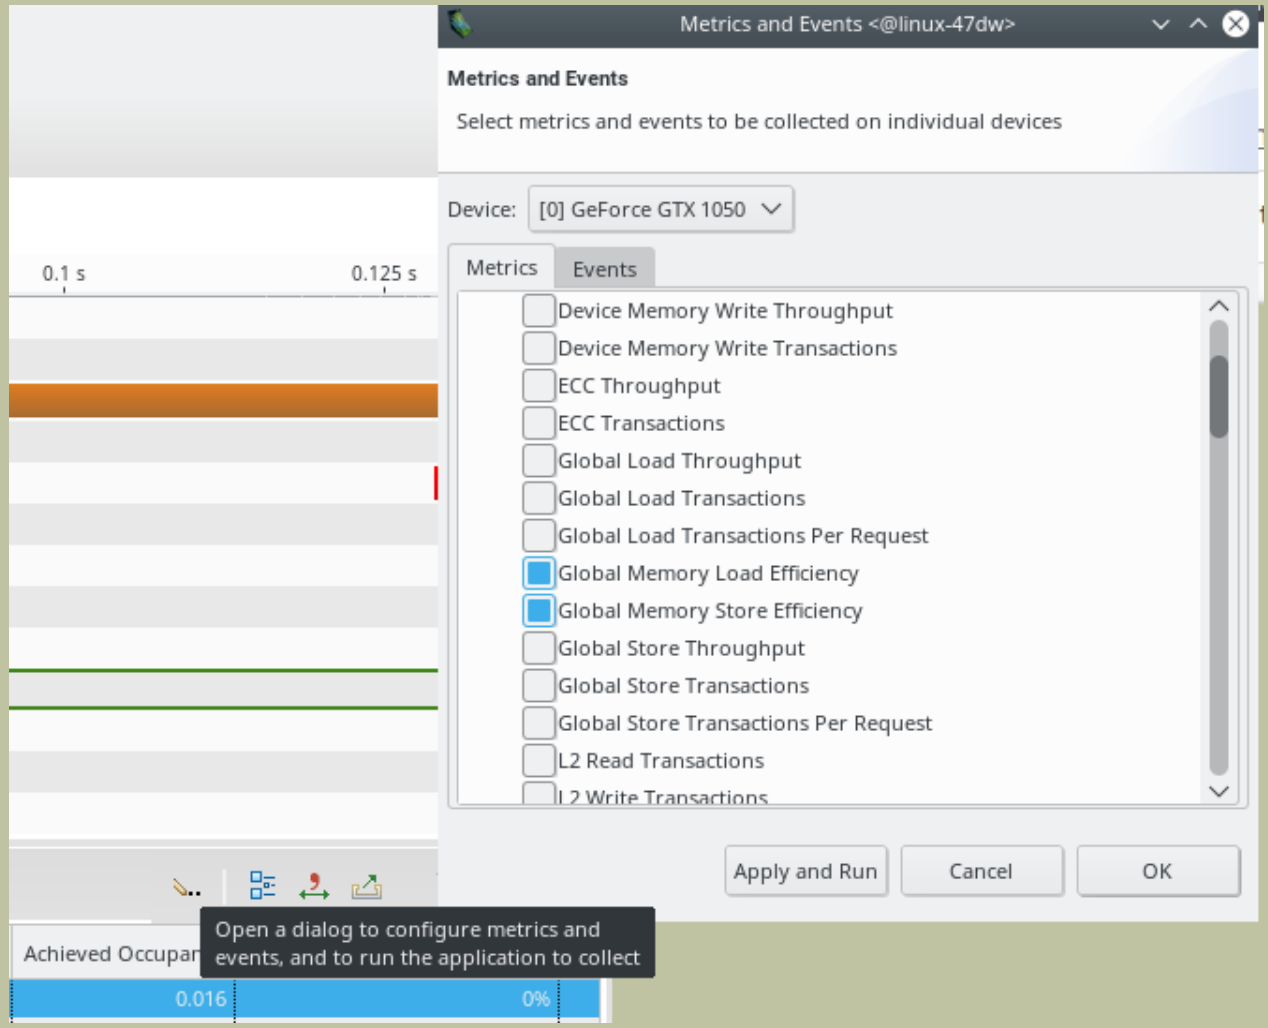
Metrics and Events <@linux-47dw>
Metrics and Events
Select metrics and events to be collected on individual devices
Device: [0] GeForce GTX 1050
Metrics
Events
Device Memory Write Throughput
Device Memory Write Transactions
ECC Throughput
ECC Transactions
Global Load Throughput
Global Load Transactions
Global Load Transactions Per Request
Global Memory Load Efficiency
Global Memory Store Efficiency
Global Store Throughput
Global Store Transactions
Global Store Transactions Per Request
L2 Read Transactions
L2 Write Transactions
Apply and Run
Cancel
OK
0.1 s
0.125 s
Open a dialog to configure metrics and events, and to run the application to collect
Achieved Occupan
0.016
0%

```
/Лекция4/lab4> ncu-ui  &
```

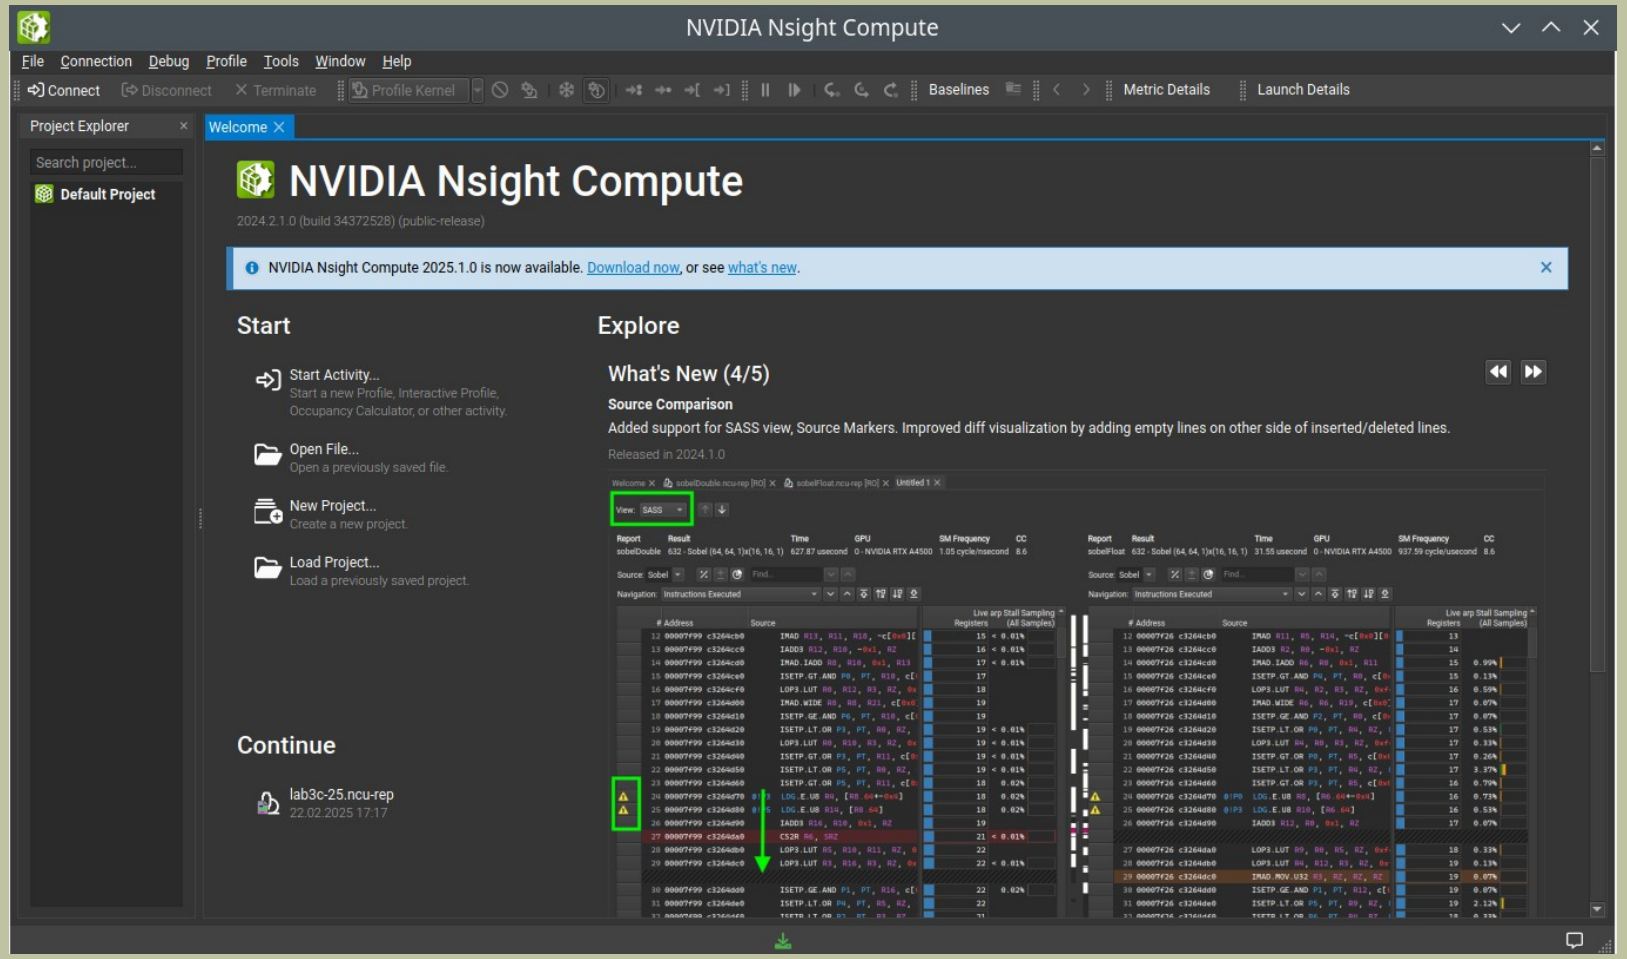

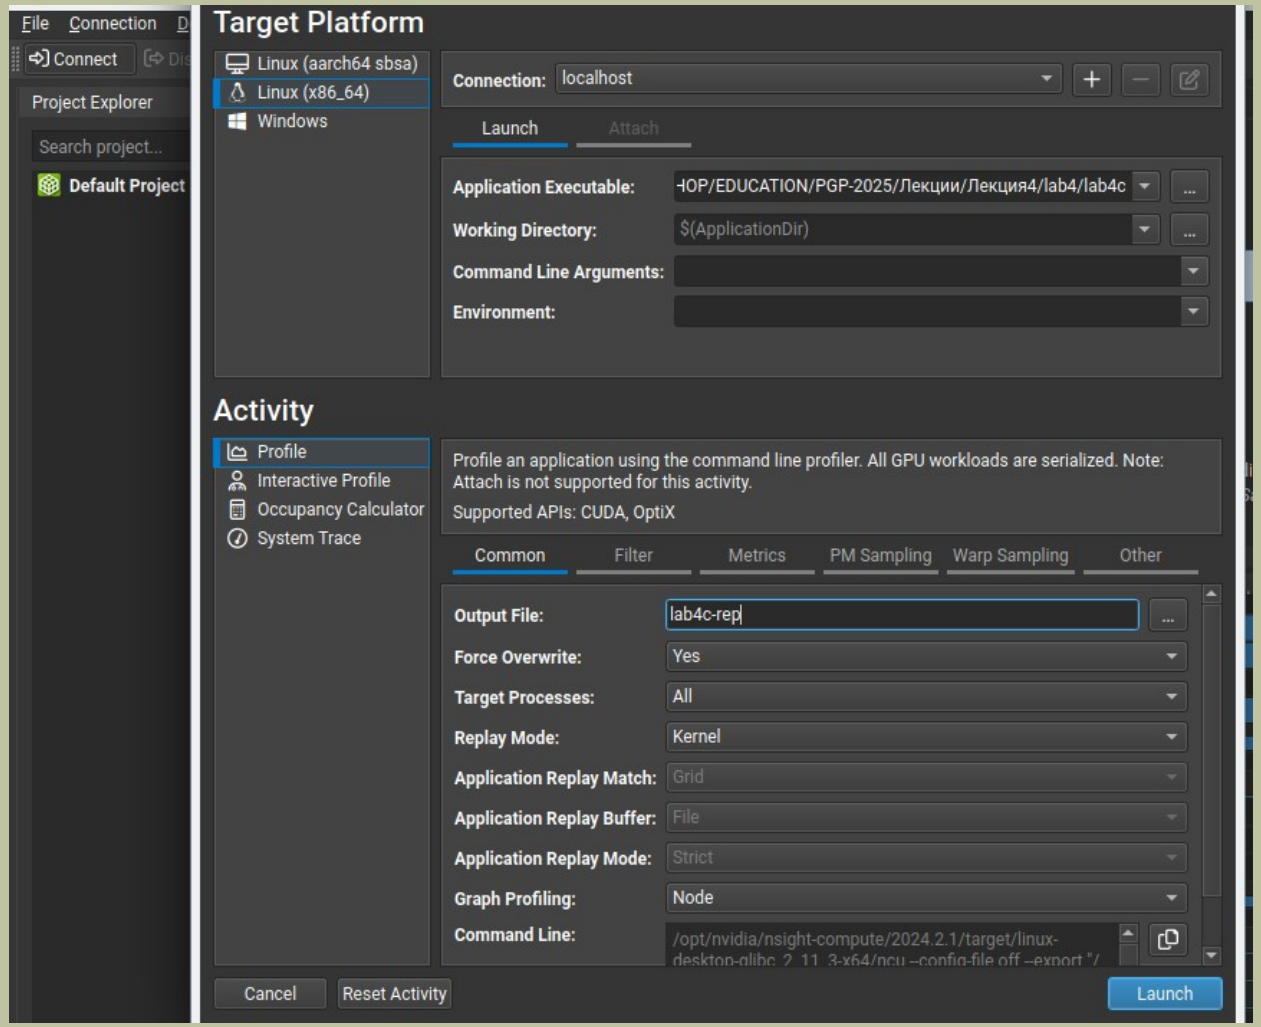
File Connection D
Connect
Project Explorer
Search project...
Default Project
Target Platform
Linux (aarch64 sbsa)
Linux (x86_64)
Windows
Connection: localhost
Launch
Attach
Application Executable: НОР/EDUCATION/PGP-2025/Лекции/Лекция4/lab4/lab4c
Working Directory: $(ApplicationDir)
Command Line Arguments:
Environment:
Activity
Profile
Interactive Profile
Occupancy Calculator
System Trace
Profile an application using the command line profiler. All GPU workloads are serialized. Note: Attach is not supported for this activity.
Supported APIs: CUDA, OptiX
Common
Filter
Metrics
PM Sampling
Warp Sampling
Other
Output File: lab4c-rep
Force Overwrite: Yes
Target Processes: All
Replay Mode: Kernel
Application Replay Match: Grid
Application Replay Buffer: File
Application Replay Mode: Strict
Graph Profiling: Node
Command Line: /opt/nvidia/nsight-compute/2024.2.1/target/linux-
Cancel
Reset Activity
Launch

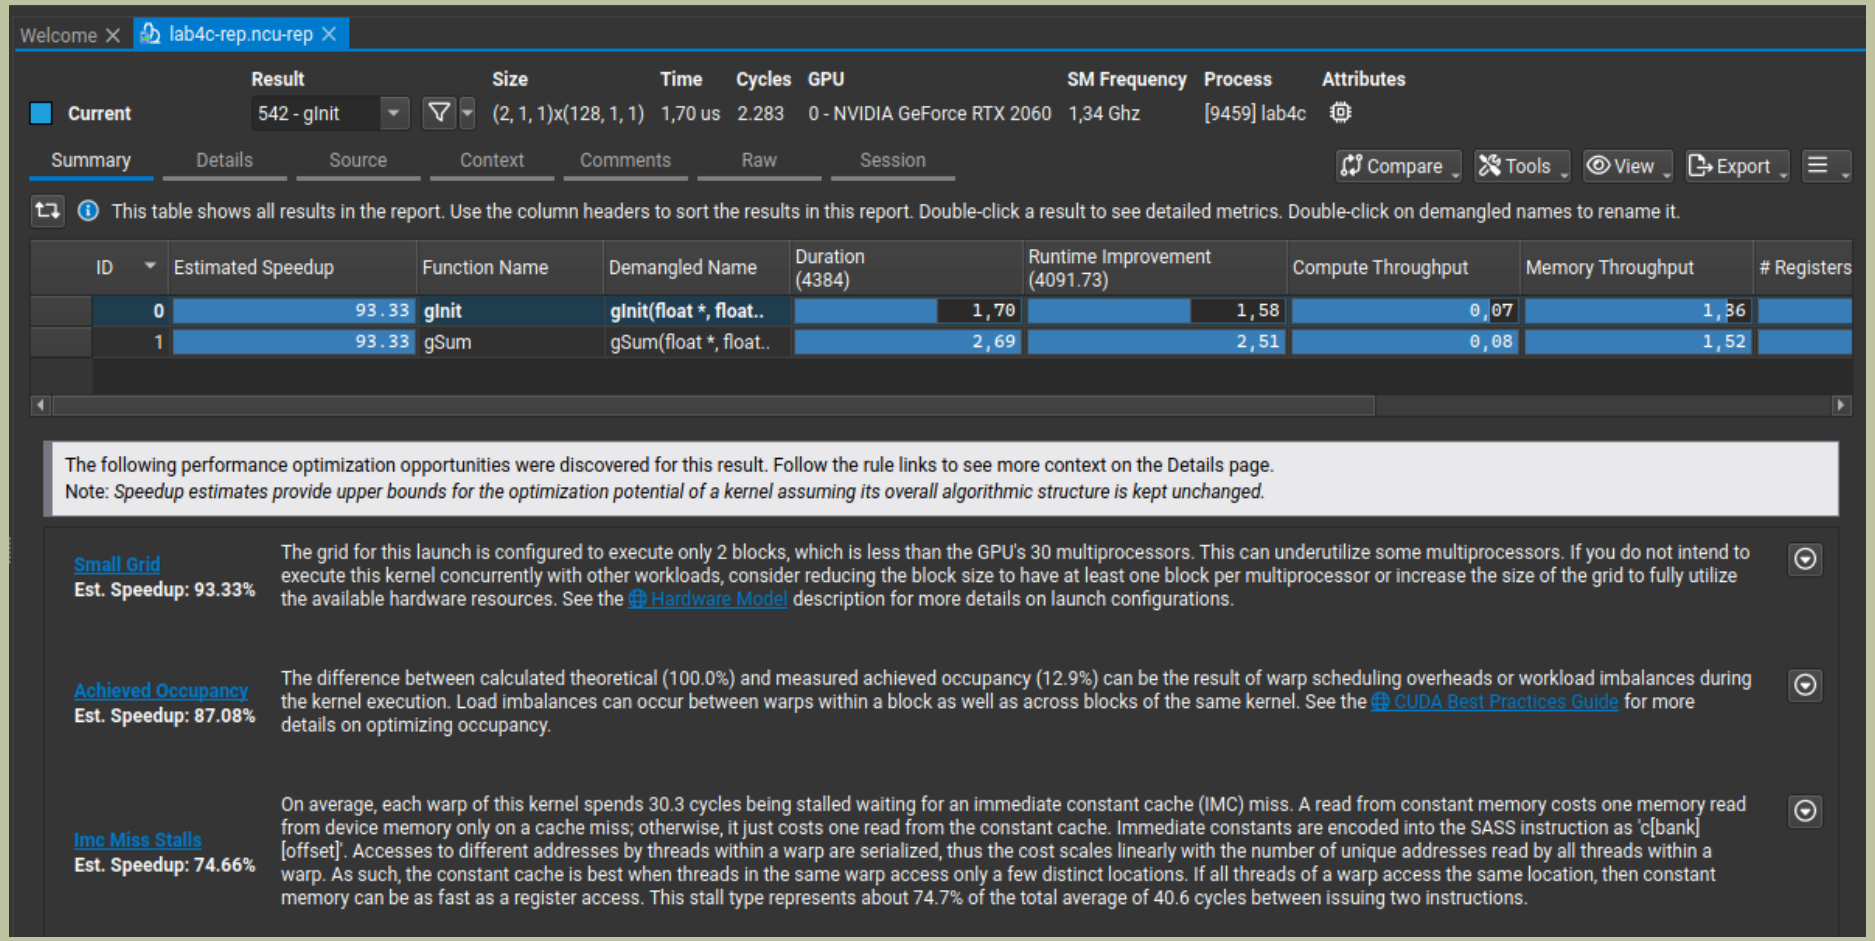
Welcome
lab4c-rep.ncu-rep
Result
Size
Time
Cycles
GPU
SM Frequency
Process
Attributes
Current
542 - gInit
(2, 1, 1)x(128, 1, 1)
1,70 us
2.283
0 - NVIDIA GeForce RTX 2060
1,34 Ghz
[9459] lab4c
Summary
Details
Source
Context
Comments
Raw
Session
Compare
Tools
View
Export
This table shows all results in the report. Use the column headers to sort the results in this report. Double-click a result to see detailed metrics. Double-click on demangled names to rename it.
ID
Estimated Speedup
Function Name
Demangled Name
Duration (4384)
Runtime Improvement (4091.73)
Compute Throughput
Memory Throughput
Registers
0
93.33
gInit
gInit(float *, float..
1,70
1,58
0,07
1,36
1
93.33
gSum
gSum(float *, float..
2,69
2,51
0,08
1,52
The following performance optimization opportunities were discovered for this result. Follow the rule links to see more context on the Details page.
Note: Speedup estimates provide upper bounds for the optimization potential of a kernel assuming its overall algorithmic structure is kept unchanged.
Small Grid
Est. Speedup: 93.33%
The grid for this launch is configured to execute only 2 blocks, which is less than the GPU's 30 multiprocessors. This can underutilize some multiprocessors. If you do not intend to execute this kernel concurrently with other workloads, consider reducing the block size to have at least one block per multiprocessor or increase the size of the grid to fully utilize the available hardware resources. See the Hardware Model description for more details on launch configurations.
Achieved Occupancy
Est. Speedup: 87.08%
The difference between calculated theoretical (100.0%) and measured achieved occupancy (12.9%) can be the result of warp scheduling overheads or workload imbalances during the kernel execution. Load imbalances can occur between warps within a block as well as across blocks of the same kernel. See the CUDA Best Practices Guide for more details on optimizing occupancy.
Imc Miss Stalls
Est. Speedup: 74.66%
On average, each warp of this kernel spends 30.3 cycles being stalled waiting for an immediate constant cache (IMC) miss. A read from constant memory costs one memory read from device memory only on a cache miss; otherwise, it just costs one read from the constant cache. Immediate constants are encoded into the SASS instruction as 'c[bank][offset]'. Accesses to different addresses by threads within a warp are serialized, thus the cost scales linearly with the number of unique addresses read by all threads within a warp. As such, the constant cache is best when threads in the same warp access only a few distinct locations. If all threads of a warp access the same location, then constant memory can be as fast as a register access. This stall type represents about 74.7% of the total average of 40.6 cycles between issuing two instructions.

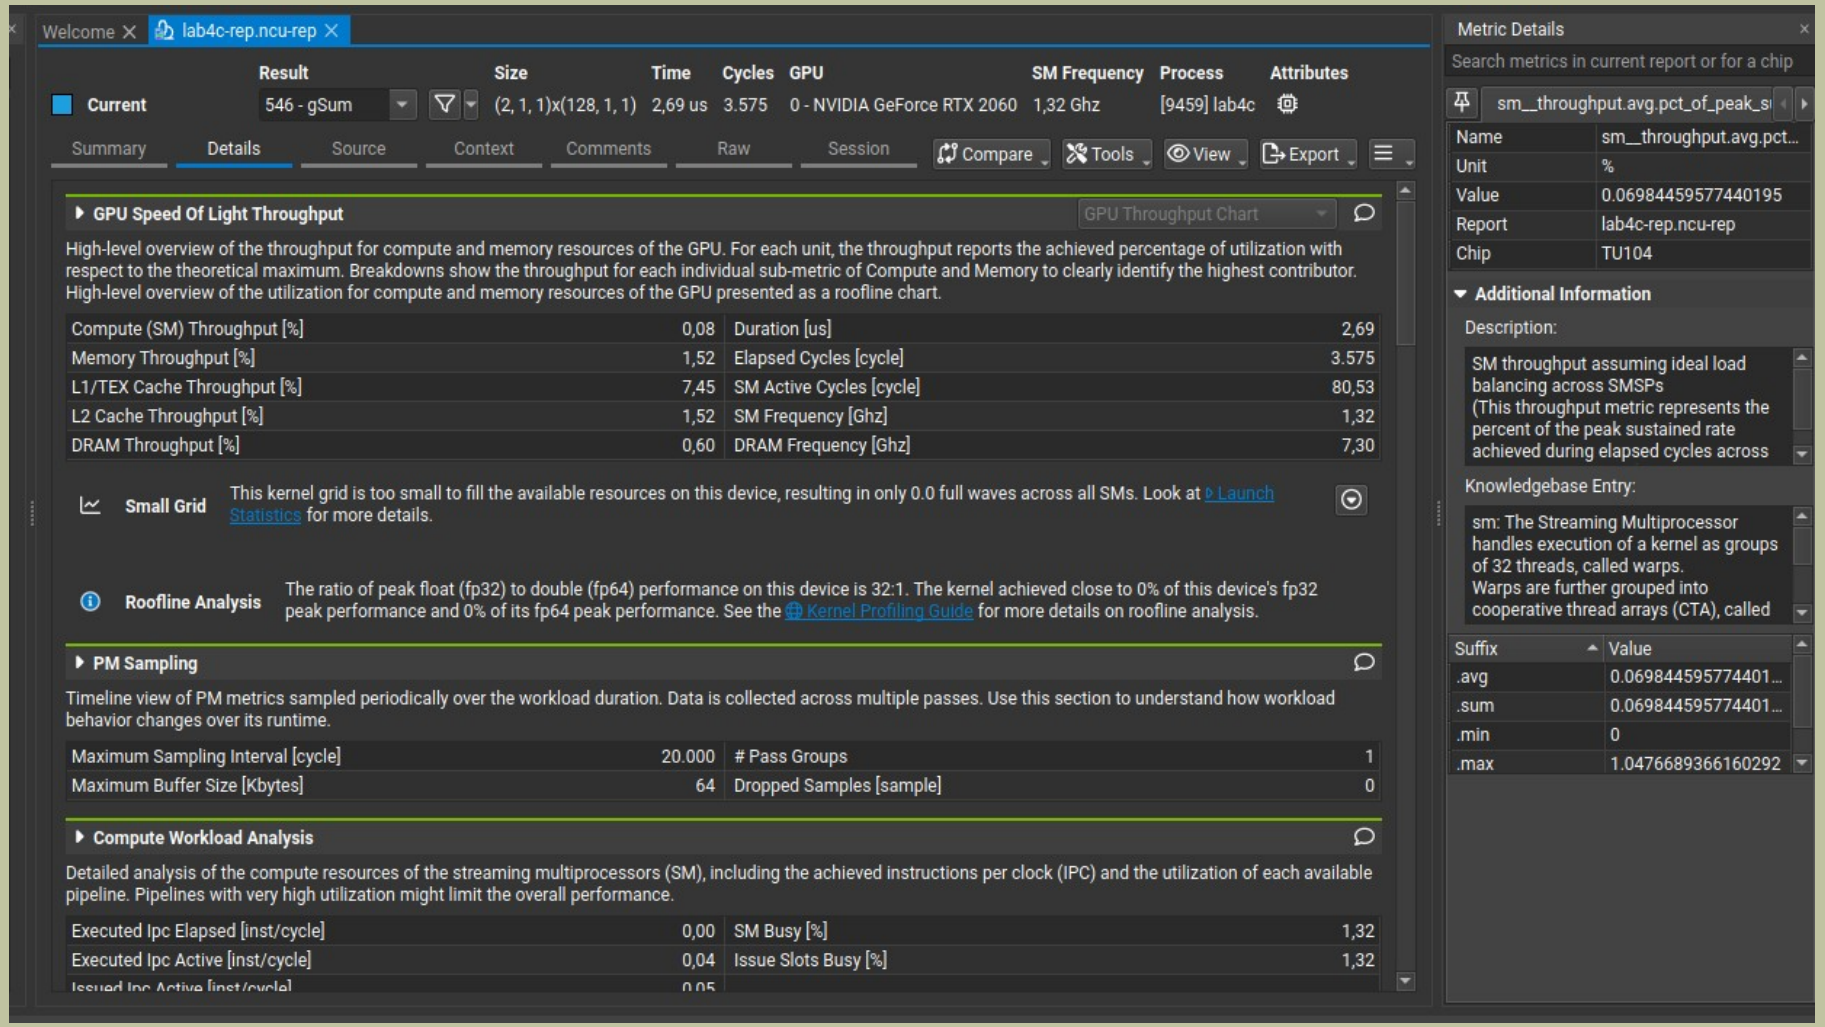
Welcome
lab4c-rep.ncu-rep
Result
Size
Time
Cycles
GPU
SM Frequency
Process
Attributes
Current
546 - gSum
(2, 1, 1)x(128, 1, 1)
2,69 us
3.575
0 - NVIDIA GeForce RTX 2060
1,32 Ghz
[9459] lab4c
Summary
Details
Source
Context
Comments
Raw
Session
Compare
Tools
View
Export
GPU Speed Of Light Throughput
GPU Throughput Chart
High-level overview of the throughput for compute and memory resources of the GPU. For each unit, the throughput reports the achieved percentage of utilization with respect to the theoretical maximum. Breakdowns show the throughput for each individual sub-metric of Compute and Memory to clearly identify the highest contributor. High-level overview of the utilization for compute and memory resources of the GPU presented as a roofline chart.
Compute (SM) Throughput [%] 0,08
Memory Throughput [%] 1,52
L1/TEX Cache Throughput [%] 7,45
L2 Cache Throughput [%] 1,52
DRAM Throughput [%] 0,60
Duration [us] 2,69
Elapsed Cycles [cycle] 3.575
SM Active Cycles [cycle] 80,53
SM Frequency [Ghz] 1,32
DRAM Frequency [Ghz] 7,30
Small Grid
This kernel grid is too small to fill the available resources on this device, resulting in only 0.0 full waves across all SMs. Look at Launch Statistics for more details.
Roofline Analysis
The ratio of peak float (fp32) to double (fp64) performance on this device is 32:1. The kernel achieved close to 0% of this device's fp32 peak performance and 0% of its fp64 peak performance. See the Kernel Profiling Guide for more details on roofline analysis.
PM Sampling
Timeline view of PM metrics sampled periodically over the workload duration. Data is collected across multiple passes. Use this section to understand how workload behavior changes over its runtime.
Maximum Sampling Interval [cycle] 20.000
Maximum Buffer Size [Kbytes] 64
Pass Groups 1
Dropped Samples [sample] 0
Compute Workload Analysis
Detailed analysis of the compute resources of the streaming multiprocessors (SM), including the achieved instructions per clock (IPC) and the utilization of each available pipeline. Pipelines with very high utilization might limit the overall performance.
Executed Ipc Elapsed [inst/cycle] 0,00
Executed Ipc Active [inst/cycle] 0,04
Issued Ipc Active [inst/cycle] 0,05
SM Busy [%] 1,32
Issue Slots Busy [%] 1,32
Metric Details
Search metrics in current report or for a chip
sm__throughput.avg.pct_of_peak_s
Name sm__throughput.avg.pct...
Unit %
Value 0.06984459577440195
Report lab4c-rep.ncu-rep
Chip TU104
Additional Information
Description:
SM throughput assuming ideal load balancing across SMSPs
(This throughput metric represents the percent of the peak sustained rate achieved during elapsed cycles across
Knowledgebase Entry:
sm: The Streaming Multiprocessor handles execution of a kernel as groups of 32 threads, called warps.
Warps are further grouped into cooperative thread arrays (CTA), called
Suffix
Value
.avg 0.069844595774401...
.sum 0.069844595774401...
.min 0
.max 1.0476689366160292

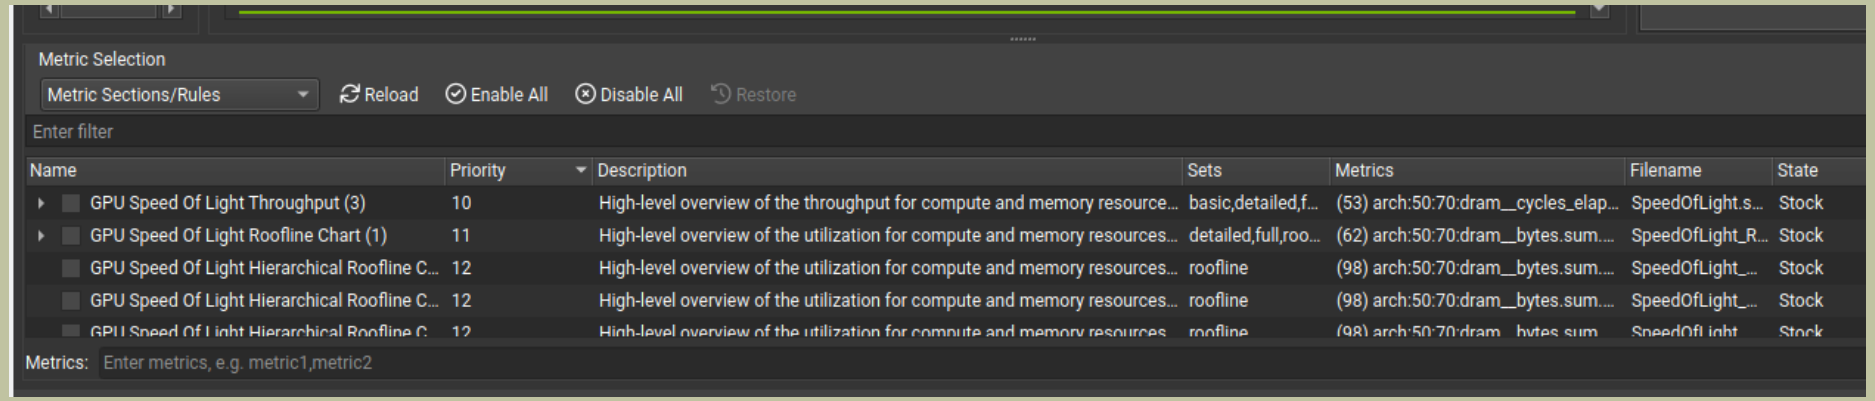
Metric Selection
Metric Sections/Rules
Reload
Enable All
Disable All
Restore
Enter filter
Name
Priority
Description
Sets
Metrics
Filename
State
GPU Speed Of Light Throughput (3)
10
High-level overview of the throughput for compute and memory resource...
basic,detailed,f...
(53) arch:50:70:dram__cycles_elap...
SpeedOfLight.s...
Stock
GPU Speed Of Light Roofline Chart (1)
11
High-level overview of the utilization for compute and memory resources...
detailed,full,roo...
(62) arch:50:70:dram__bytes.sum....
SpeedOfLight_R...
Stock
GPU Speed Of Light Hierarchical Roofline C...
12
High-level overview of the utilization for compute and memory resources...
roofline
(98) arch:50:70:dram__bytes.sum....
SpeedOfLight_...
Stock
GPU Speed Of Light Hierarchical Roofline C...
12
High-level overview of the utilization for compute and memory resources...
roofline
(98) arch:50:70:dram__bytes.sum....
SpeedOfLight_...
Stock
Metrics:
Enter metrics, e.g. metric1,metric2

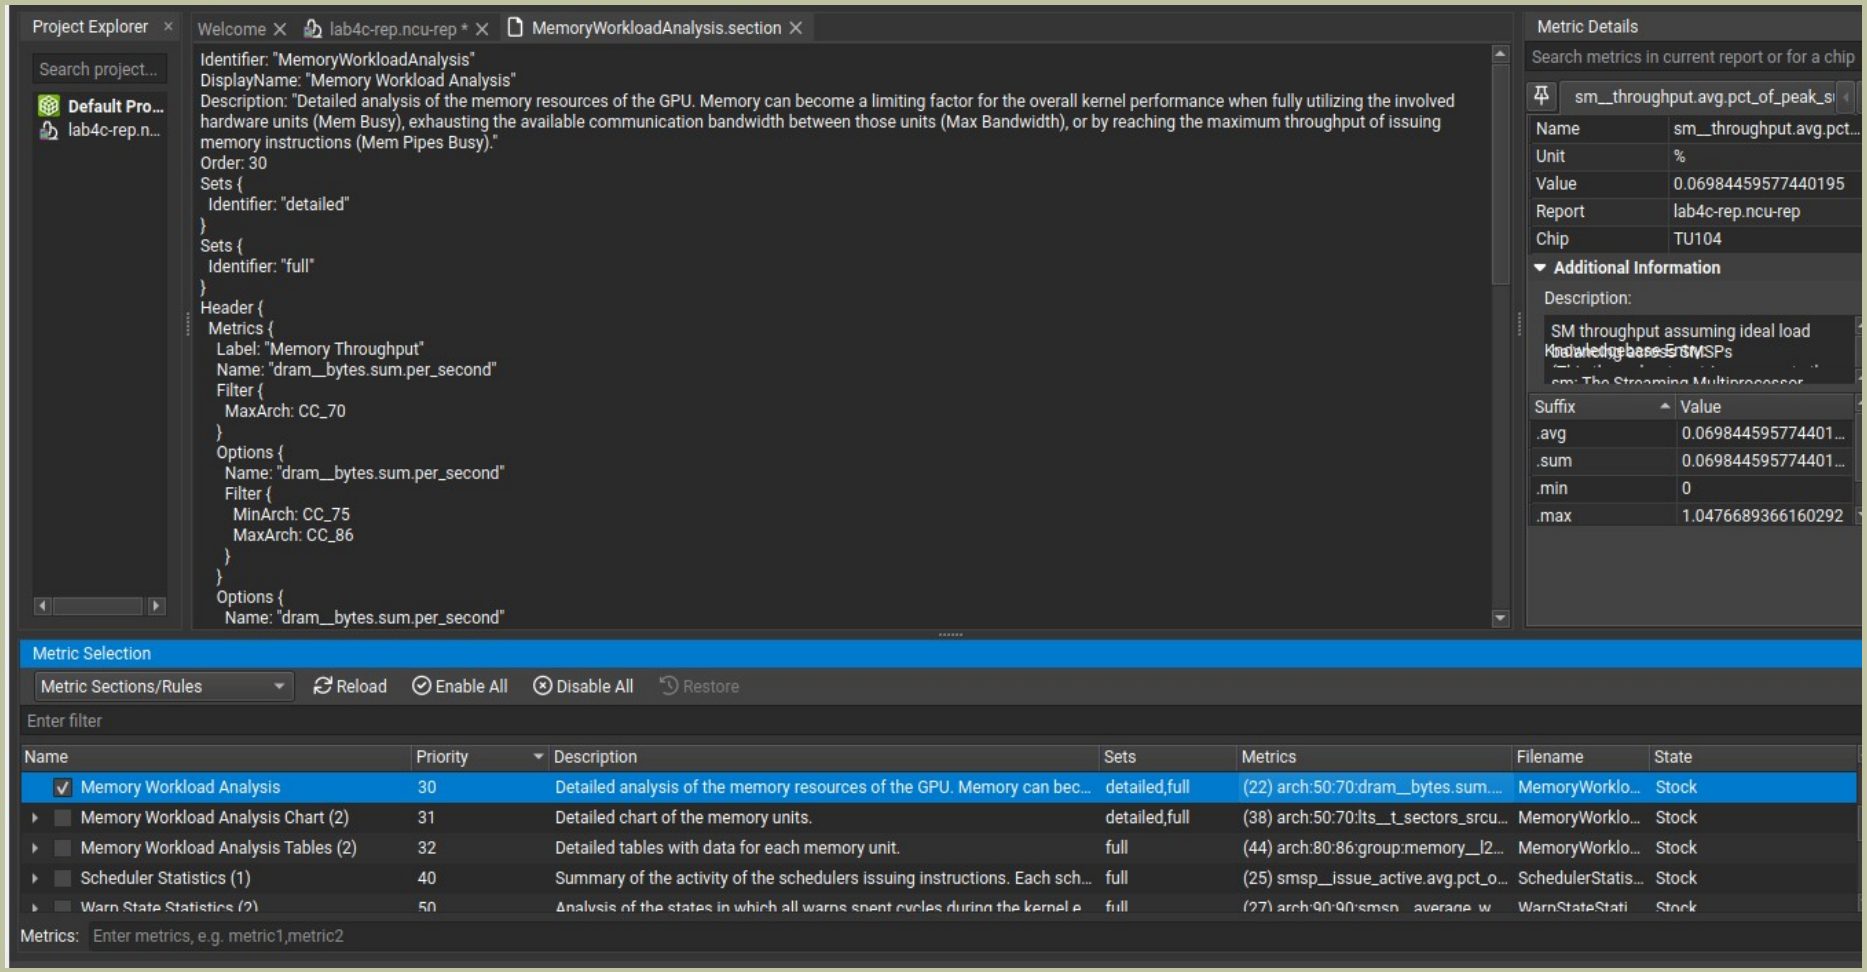
Project Explorer
Search project...
Default Pro...
lab4c-rep.n...
Welcome
lab4c-rep.ncu-rep *
MemoryWorkloadAnalysis.section
Identifier: "MemoryWorkloadAnalysis"
DisplayName: "Memory Workload Analysis"
Description: "Detailed analysis of the memory resources of the GPU. Memory can become a limiting factor for the overall kernel performance when fully utilizing the involved hardware units (Mem Busy), exhausting the available communication bandwidth between those units (Max Bandwidth), or by reaching the maximum throughput of issuing memory instructions (Mem Pipes Busy)."
Order: 30
Sets {
Identifier: "detailed"
}
Sets {
Identifier: "full"
}
Header {
Metrics {
Label: "Memory Throughput"
Name: "dram__bytes.sum.per_second"
Filter {
MaxArch: CC_70
}
Options {
Name: "dram__bytes.sum.per_second"
Filter {
MinArch: CC_75
MaxArch: CC_86
}
}
Options {
Name: "dram__bytes.sum.per_second"
Metric Details
Search metrics in current report or for a chip
sm__throughput.avg.pct_of_peak_s
Name
sm__throughput.avg.pct...
Unit
%
Value
0.06984459577440195
Report
lab4c-rep.ncu-rep
Chip
TU104
Additional Information
Description:
SM throughput assuming ideal load
Suffix
Value
.avg
0.069844595774401...
.sum
0.069844595774401...
.min
0
.max
1.0476689366160292
Metric Selection
Metric Sections/Rules
Reload
Enable All
Disable All
Restore
Enter filter
Name
Priority
Description
Sets
Metrics
Filename
State
Memory Workload Analysis
30
Detailed analysis of the memory resources of the GPU. Memory can bec...
detailed,full
(22) arch:50:70:dram__bytes.sum....
MemoryWorklo...
Stock
Memory Workload Analysis Chart (2)
31
Detailed chart of the memory units.
detailed,full
(38) arch:50:70:lts__t_sectors_srcu...
MemoryWorklo...
Stock
Memory Workload Analysis Tables (2)
32
Detailed tables with data for each memory unit.
full
(44) arch:80:86:group:memory__l2...
MemoryWorklo...
Stock
Scheduler Statistics (1)
40
Summary of the activity of the schedulers issuing instructions. Each sch...
full
(25) smsp__issue_active.avg.pct_o...
SchedulerStatis...
Stock
Metrics:
Enter metrics, e.g. metric1,metric2